

# CS330: Operating Systems

Condition variables, Concurrency bugs

# Condition variables

*pthread\_cond\_wait(pthread\_cond\_t \*cond, pthread\_mutex\_t \*mutex):*

Atomically releases the mutex and waits on a condition variable. Resumes execution holding the lock when pthread\_cond\_signal( ) is invoked.

Important: *The caller should perform the condition check after wakeup*

*pthread\_cond\_signal(pthread\_cond\_t \*cond)*

Wakes up a waiting thread on condition *cond*, Ideally called holding the mutex.

# Example usage

```
cond_t *C; lock_t *L;  bool condition;
```

```
void T1()  
{  
    while(1){  
        lock(L);  
        while(condition != true)  
            cond_wait(C, L);  
        unlock(L);  
        process();  
    }  
}
```

```
void T2()  
{  
    while(1){  
        condition = false;  
        process();  
        lock(L);  
        condition = true;  
        cond_signal(C);  
        unlock(L);  
    }  
}
```

# Example usage

```
cond_t *C; lock_t *L;  bool condition;
```

```
void T1()
```

- Why explicit condition check is required (in the waiting thread)?
- Why lock must be held while invoking cond\_signal( )?

```
while(condition != true)
    cond_wait(C, L);
unlock(L);
process();
}
```

```
void T1()
```

```
lock(L);
condition = true;
cond_signal(C);
unlock(L);
}
```

# Example usage

```
cond_t *C; lock_t *L;  bool condition;
```

- Why explicit condition check is required (in the waiting thread)?
- Some implementation on multicore may wake up more than one thread (cause spurious wakeups). For more information, please refer the man page [https://linux.die.net/man/3/pthread\\_cond\\_wait](https://linux.die.net/man/3/pthread_cond_wait).
- Why lock must be held while invoking cond\_signal( )?

```
unlock(L);  
process();  
}  
}
```

```
condition = true;  
cond_signal(C);  
unlock(L);  
}  
}
```

# Example usage

```
cond_t *C; lock_t *L;  bool condition;
```

- Why explicit condition check is required (in the waiting thread)?
- Some implementation on multicore may wake up more than one thread (cause spurious wakeups). For more information, please refer the man page [https://linux.die.net/man/3/pthread\\_cond\\_wait](https://linux.die.net/man/3/pthread_cond_wait) .
- Why lock must be held while invoking `cond_signal( )`?
- The waiting thread may wait indefinitely when the signaling thread executes `cond_signal( )` before the waiter invokes `cond_wait( )`

# Semaphore using condition variables (ostep-31.17)

```
struct sem_t {  
    int value;  
    lock_t lock;  
    cond_t cond;  
};  
  
void sem_wait( sem_t *S)  
{  
    lock(&S → lock);  
    while( S → value <= 0)  
        cond_wait(&S → cond, &S → lock);  
    S → value --;  
    unlock(&S → lock);  
}
```

```
void sem_init(sem_t *S, int val)  
{  
    S → value = val;  
    cond_init(&S → cond);  
    lock_init(&S → lock);  
}  
  
void sem_post( sem_t *S)  
{  
    lock(&S → lock); S → value ++;  
    cond_signal(&S → cond);  
    unlock(&S → lock);  
}
```

# Concurrency bugs - atomicity issues

```
char *ptr; // Allocated before use
```

```
void T1()
```

```
{
```

```
...
```

```
strcpy(ptr, "hello world!");
```

```
...
```

```
}
```

```
void T2()
```

```
{
```

```
...
```

```
if(some_condition)
```

```
free(ptr);
```

```
...
```

```
}
```

- This code is buggy. What is the issue?



# Concurrency bugs - atomicity issues

```
char *ptr; // Allocated before use

void T1()
{
    ...
    strcpy(ptr, "hello world!");
    ...
}

void T2()
{
    ...
    if(some_condition)
        free(ptr);
    ...
}
```

- This code is buggy. What is the issue?
- T2 can free the pointer before T1 uses it.
- How to fix it?

# Concurrency bugs - atomicity issues

```
char *ptr; // Allocated before use
```

```
void T1()
```

```
{
```

```
...
```

```
if(ptr) strcpy(ptr, "hello world!");
```

```
...
```

```
}
```

```
void T2()
```

```
{
```

```
...
```

```
if(some_condition)
```

```
free(ptr);
```

```
...
```

```
}
```

- Does the above fix (checking ptr in T1) work?

# Concurrency bugs - atomicity issues

```
char *ptr; // Allocated before use

void T1()
{
    ...
    if(ptr) strcpy(ptr, "hello world!");
    ...
}

void T2()
{
    ...
    if(some_condition)
        free(ptr);
    ...
}
```

- Does the above fix (checking ptr in T1) work?
- Not really. Consider the following order of execution:
- T1: "if(ptr)" T2: "free(ptr)" T1: "strcpy" Result: Segfault

# Concurrency bugs - ordering issues

```
1.  bool pending;  
2.  void T1()  
3.  {  
4.      pending = true;  
5.      do_large_processing();  
6.      while (pending);  
7.  }
```

```
1.  void T2()  
2.  {  
3.      do_some_processing();  
4.      pending = false;  
5.      some_other_processing();  
6.  }
```

- This code works with the assumption that line#4 of T2 is executed after line#4 of T1
- If this ordering is violated, T1 is stuck in the while loop

# Concurrency bugs - deadlocks

```
struct acc_t{
    lock_t *L;
    id_t acc_no;
    long balance;
}

void txn_transfer( acc_t *src,
                  acc_t *dst, long amount)
{
    lock(src → L); lock(dst → L);
    check_and_transfer(src, amount);
    unlock(dst → L); unlock(src → L);
}
```

- Consider a simple transfer transaction in a bank
- Where is the deadlock?


# Concurrency bugs - deadlocks

```
struct acc_t{
    lock_t *L;
    id_t acc_no;
    long balance;
}

void txn_transfer( acc_t *src,
                  acc_t *dst, long amount)
{
    lock(src → L); lock(dst → L);
    check_and_transfer(src, amount);
    unlock(dst → L); unlock(src → L);
}
```

- Consider a simple transfer transaction in a bank
- Where is the deadlock?
- T1: txn\_transfer(iitk, cse, 10000)
  - lock (iitk), lock (cse)
- T2: txn\_transfer(cse, iitk, 5000)
  - lock (cse), lock(iitk)

# Dining philosophers



```
atomic_t forks[5];
Philosopher( int id)
{
    while (1) {
        think( );
        acquire(forks[id]);
        acquire(forks[(id+1) % 5]);
        eat( );
        release( forks[(id+1) % 5]);
        release(forks[id]);
    }
}
```

# Conditions for deadlock

- Mutual exclusion: exclusive control of resources (e.g, thread holding lock)
- Hold-and-wait: hold one resource and wait for other
- No resource preemption: Resources can not be forcibly removed from threads holding them
- Circular wait: A cycle of threads requesting locks held by others. Specifically, a cycle in the directed graph  $G(V, E)$  where  $V$  is the set of processes and  $(v_1, v_2) \in E$  if  $v_1$  is waiting for a lock held by  $v_2$

All of the above conditions should be satisfied for a deadlock to occur



# Solutions for deadlocks

- Remove mutual exclusion: lock free data structures
- Either acquire all resources or no resource
  - trylock(lock) APIs can be used (e.g., pthread\_mutex\_trylock( ))
- Careful scheduling: Avoid scheduling threads such that no deadlock occur
- Most commonly used technique is to avoid circular wait. This can be achieved by ordering the resources and acquiring them in a particular order from all the threads.

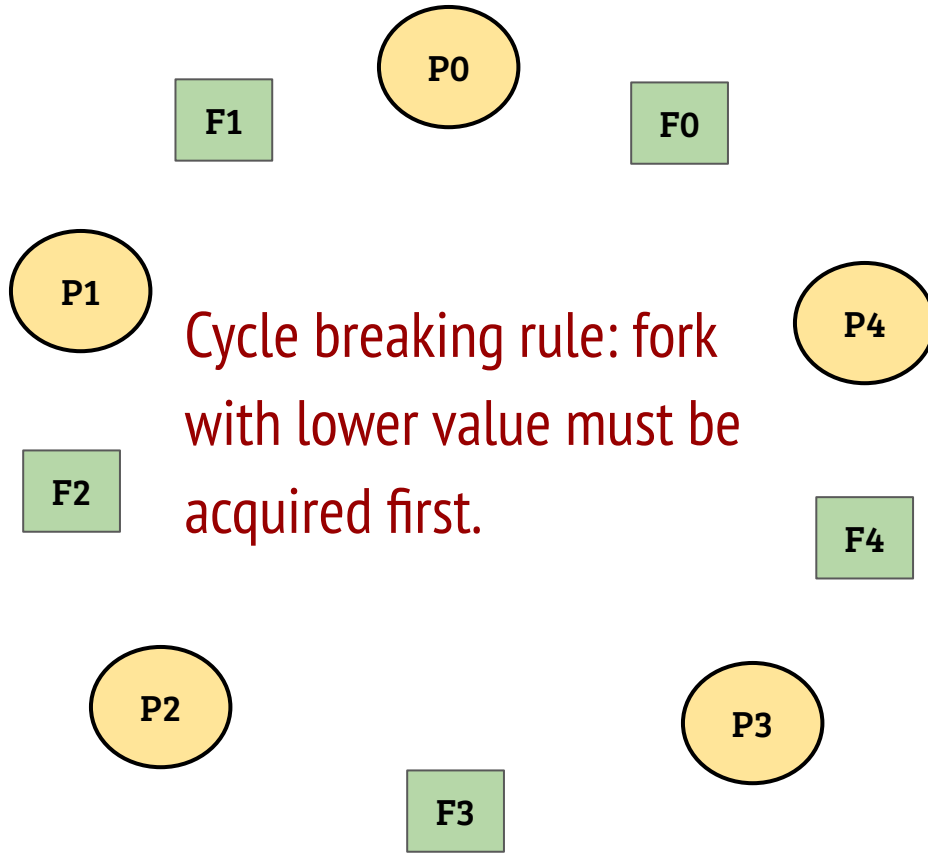
# Concurrency bugs - avoiding deadlocks

```
struct acc_t{
    lock_t *L;
    id_t acc_no;
    long balance;
}

void txn_transfer( acc_t *src,
                  acc_t *dst, long amount)
{
    lock(src → L); lock(dst → L);
    check_and_transfer(src, amount);
    unlock(dst → L); unlock(src → L);
}
```

- Deadlock in a simple transfer transaction in a bank
- While acquiring locks, first acquire the lock for the account with lower “acc\_no” value
- Account number comparison performed before acquiring the lock

# Dining philosophers: breaking the deadlock



```
atomic_t forks[5];
Philosopher( int id)
{
    while (1) {
        if(id == 4){
            acquire(0);
            acquire(4);
        }else{
            acquire(forks[id]);
            acquire(forks[id+1]);
        }
        ...
    }
}
```