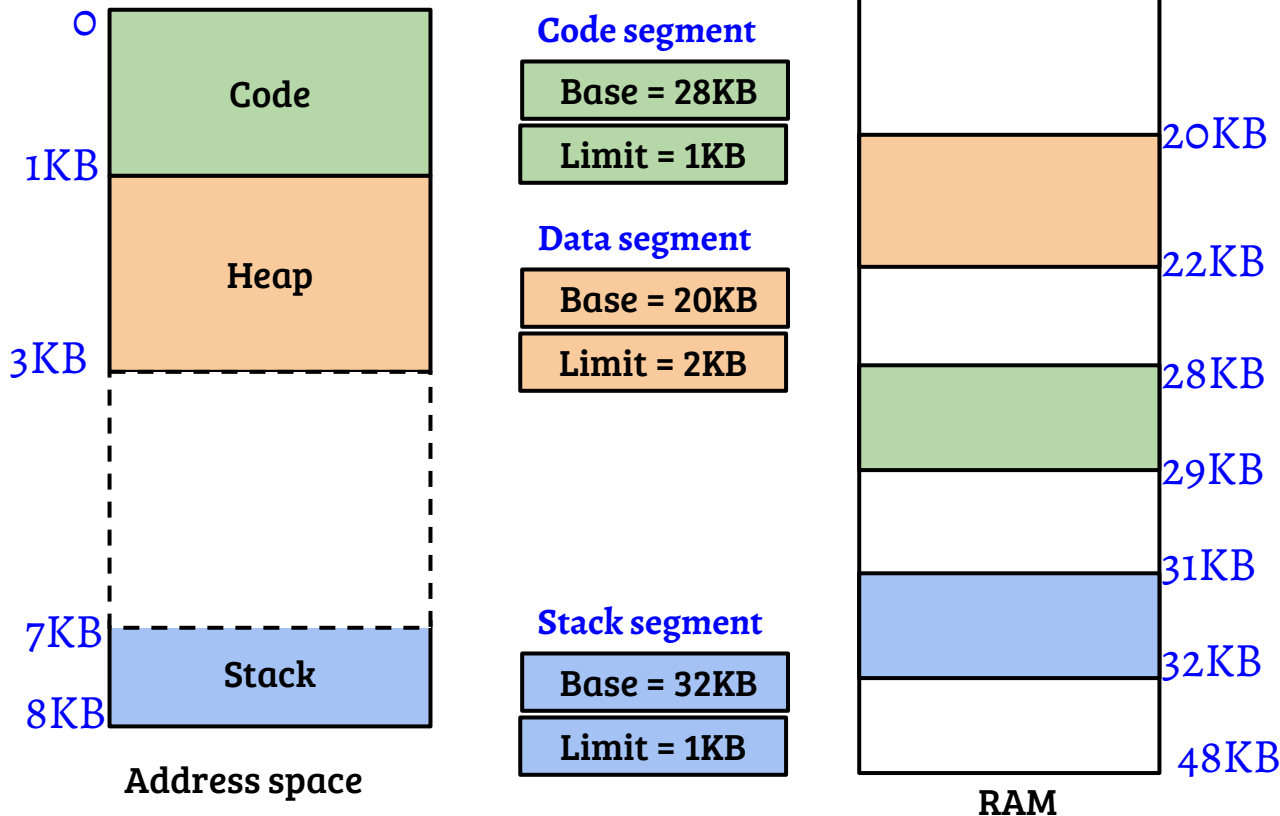


CS330: Operating Systems


Virtual memory: Paging

Recap: Segmentation



- Extension of the scheme for translation and address space granularity
- Base-limit register pairs per segment

Recap: Segmentation in reality



- Descriptor table register (DTR) is used to access the descriptor table
- # of descriptors depends on architecture
- Separate descriptors used for user and kernel mode

Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping


Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping
- The idea of paging
 - Partition the address space into fixed sized blocks (call it page)
 - Physical memory partitioned in a similar way (call it page frame)


Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping
- The idea of paging
 - Partition the address space into fixed sized blocks (call it pages)
 - Physical memory partitioned in a similar way (call it page frames)
 - OS creates a mapping between *page* to *page frame*
 - H/W uses the mapping to translate VA to PA


Paging example (pages)




- Virtual address size = 32KB, Page size = 256 bytes
- Address length = 15 bits {0x0 - 0x7FFF}
- # of pages = 128




Paging example (pages)



- Virtual address size = 32KB, Page size = 256 bytes
- Address length = 15 bits {0x0 - 0x7FFF}
- # of pages = 128




Virtual address





- Example: For Virtual address *0x0510*, Page number = 5, offset = 16


Paging example (page frames)



- Physical address size = 64KB
- Address length = 16 bits {0x0 - 0xFFFF}
- # of page frames = 256



DRAM



Process address
space


Paging example (page frames)

0	Page 0
256	Page 1
512	Page 2
768	Page 3
1024	

Page 125
Page 126
Page 127

Process address space

- Physical address size = 64KB
- Address length = 16 bits {0x0 - 0xFFFF}
- # of page frames = 256



Physical address

0	PFN 0	256
	PFN 1	512
	PFN 2	768
	PFN 3	1024

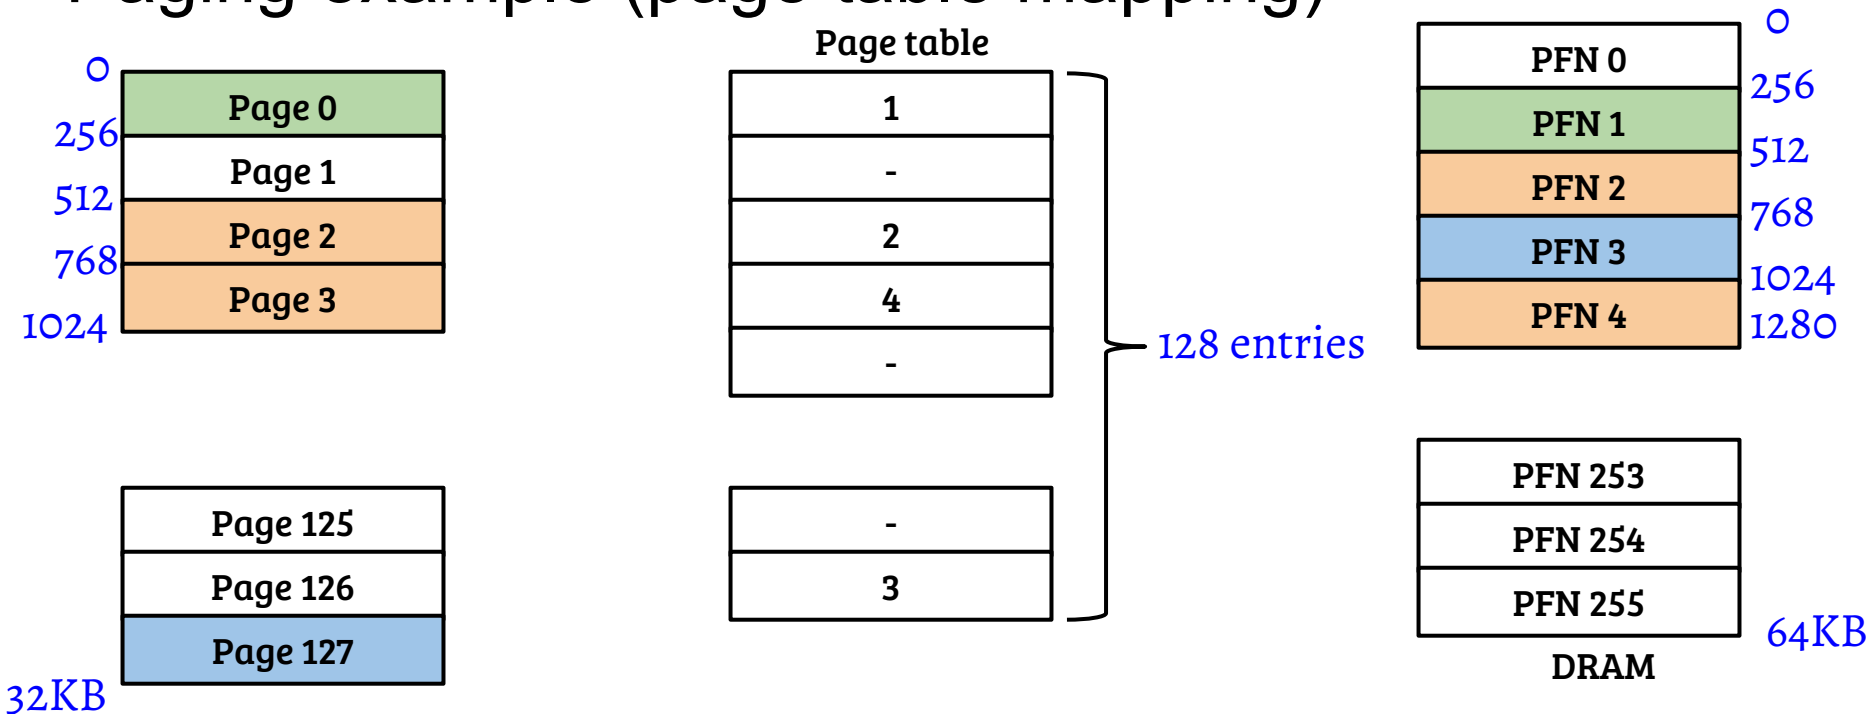
PFN 253
PFN 254
PFN 255

64KB

DRAM



- Example: For physical address *0x1F51*, PFN = 31, offset = 81

Paging example (page table mapping)





- Each entry in page table is called page table entry (PTE)
- Example mapping: page 0 \Rightarrow PFN 1, page 2 \Rightarrow PFN 2 and so on

Paging example (page table walk)




Process address
space

```
PTW (vaddr V, PTable P)
// Input: Virtual address, Page table
// Returns physical address
{
    Entry = P[V >> 8];
    if (Entry.present)
        return (Entry.PFN << 8) + (V & 0xFF);
    Raise PageFault;
}
```



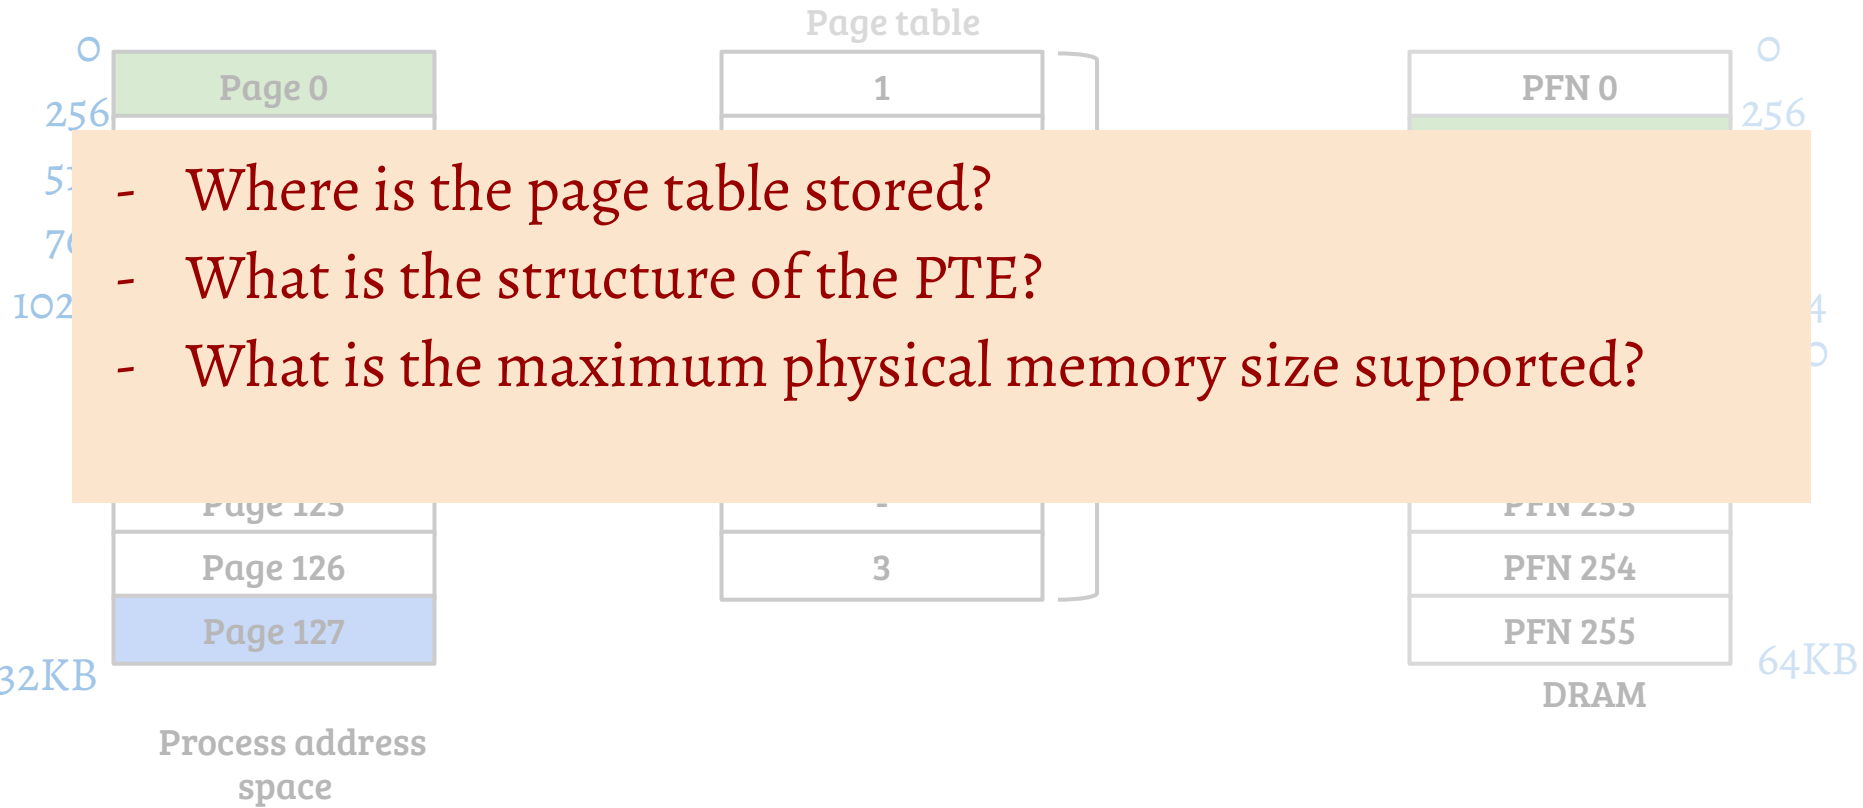
DRAM

Paging example (example translation)

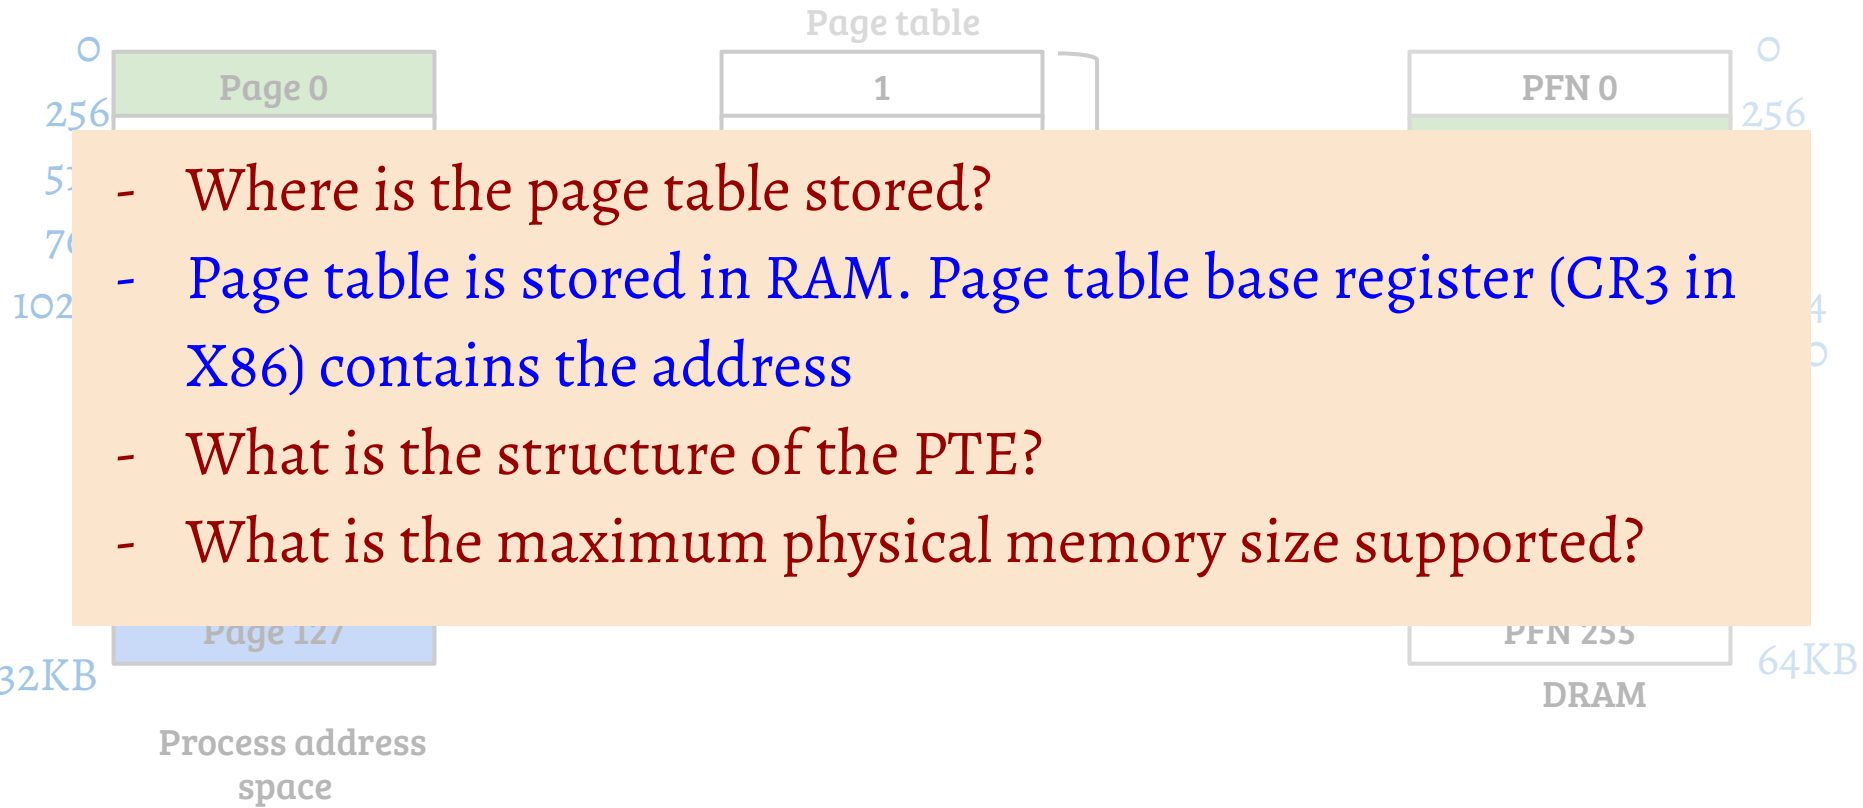


- Process address space
- Virtual address 0x10 translates to physical address 0x110
 - Virtual address 0x7FF0 translates to physical address 0x3F0

Paging example (page table walk)



Paging example (page table walk)



Paging example (structure of an example PTE)



- PFN occupies a significant portion of PTE entry (8 bits in this example)

P

Present bit, 1 \Rightarrow entry is valid

W

Write bit, 1 \Rightarrow Write allowed

S

Privilege bit, 0 \Rightarrow only kernel mode access is allowed

A

Accessed bit, 1 \Rightarrow Address accessed (set by H/W during walk)

D

Dirty bit, 1 \Rightarrow Address written (set by H/W during walk)

X

Execute bit, 1 \Rightarrow Instruction fetch allowed for this page

Reserved/unused bits

Paging example (Page table entries)

0	Page 0
256	Page 1
512	Page 2
768	Page 3
1024	

Page table	
0x125	
0x0	
0x207	
0x407	
0x0	

0	PFN 0
256	PFN 1
512	PFN 2
768	PFN 3
1024	PFN 4
1280	

Page 125
Page 126
Page 127

0x0
0x307

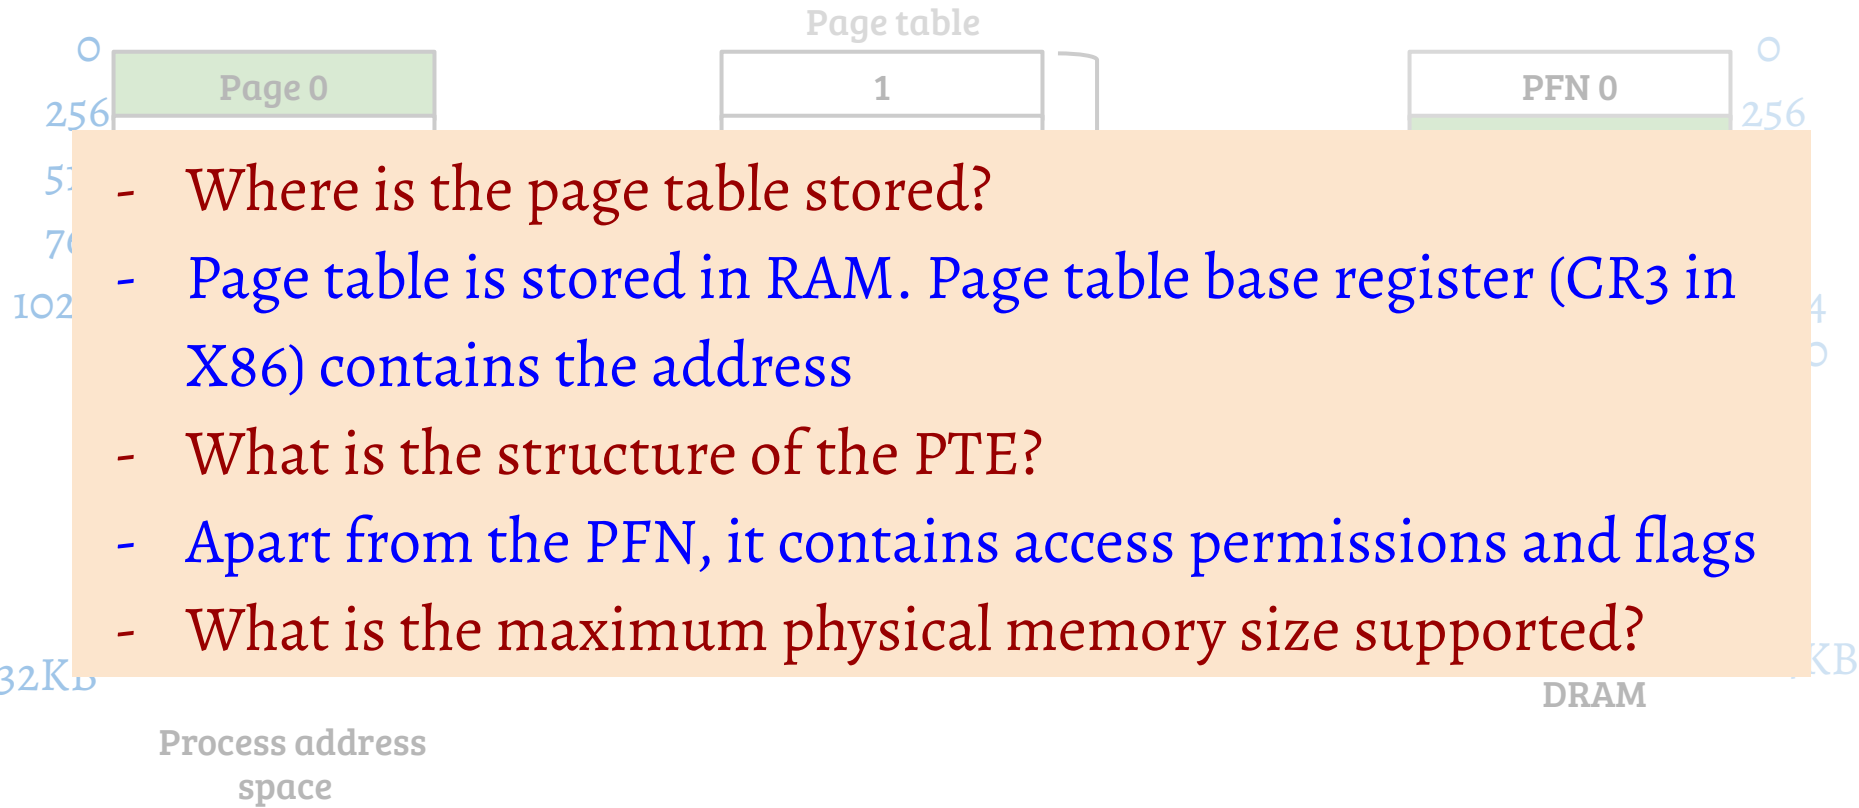
PFN 253
PFN 254
PFN 255

64KB

Process address
space

- Code: Page 0 (Read and Execute)
- Data: Page 2 and Page 3 (Read and Write)
- Stack: Page 127 (Read and Write)

Paging example (page table walk)



Paging example (page table walk)

- Where is the page table stored?
- Page table is stored in RAM. Page table base register (CR3 in X86) contains the address
- What is the structure of the PTE?
- Apart from the PFN, it contains access permissions and flags
- What is the maximum physical memory size supported?
- For this example, 8-bits can be used to specify 256 page frames. Maximum RAM size = $256 * 256 = 64\text{KB}$

Paging: one level of page table may not be feasible!

- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?


Paging: one level of page table may not be feasible!

- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?
- Large page size results in *internal fragmentation*
- Assuming page size = 4KB, How many entries are required in a one-level paging system?

Paging: one level of page table may not be feasible!

- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?
- Large page size results in *internal fragmentation*
- Assuming page size = 4KB, How many entries are required in a one-level paging system? (2^{20} entries)
- Not possible to hold 2^{20} entries in a single page
- Therefore, multi-level page tables are used in modern systems

Two-level page tables (32-bit virtual address)



- Two-level page table
- Level-1 page table contains entries pointing to Level-2 page table structures
- Level-2 entry contains PFN along with flags