

# CS330: Operating Systems

Software locks, Semaphore

# Recap: Spinlocks with hardware support

- Architectural support for atomic operations like atomic exchange, compare-and-swap, LL-SC and atomic add can be used to build spinlocks
- Ticket spinlocks provide fairness in locking, example implementation with atomic-add
- Outstanding issues: Blocking locks (will come back after semaphores)

Today's lecture: Software-only locks, Semaphores

# Buggy #1

```
int flag[2] = {0,0};  
void lock (int id)  /*id = 0 or 1 */  
{  
    while(flag[id ^ 1]); // ^  $\rightarrow$  XOR  
    flag[id] = 1;  
}  
void unlock (int id)  
{  
    flag[id] = 0;  
}
```

- Solution for two threads,  $T_0$  and  $T_1$  with id 0 and 1, respectively
- We have seen that this solution does not work, Why?

# Buggy #1

```
int flag[2] = {0,0};  
void lock (int id)  /*id = 0 or 1 */  
{  
    while(flag[id ^ 1]); // ^  $\rightarrow$  XOR  
    flag[id] = 1;  
}  
void unlock (int id)  
{  
    flag[id] = 0;  
}
```

- Solution for two threads,  $T_0$  and  $T_1$  with id 0 and 1, respectively
- We have seen that this solution does not work, Why?
- Both threads can acquire the lock as “while condition check” and “setting the flag” is non-atomic

# Buggy #2

```
int flag[2] = {0,0};
```

```
void lock (int id)  /*id = 0 or 1 */ - Does this solution work?
```

```
{
```

```
    flag[id] = 1;
```

```
    while(flag[id ^ 1]); // ^ → XOR
```

```
}
```

```
void unlock (int id)
```

```
{
```

```
    flag[id] = 0;
```

```
}
```

# Buggy #2

```
int flag[2] = {0,0};  
void lock (int id)  /*id = 0 or 1 */  
{  
    flag[id] = 1;  
    while(flag[id ^ 1]); // ^ → XOR  
}  
void unlock (int id)  
{  
    flag[id] = 0;  
}
```

- Does this solution work?
- No, as this can lead to a deadlock (flag[0] = flag[1] = 1) In other words the “progress” requirement is not met
- Progress: If no one has acquired the lock and there are contending threads, one of the threads must acquire the lock within a finite time

# Buggy #3

```
int turn = 0;
```

```
void lock (int id)  /*id = 0 or 1 */
```

```
{
```

```
    while(turn == id ^ 1);
```

```
}
```

```
void unlock (int id)
```

```
{
```

```
    turn = id ^ 1;
```

```
}
```

- Assuming  $T_0$  invokes lock( ) first, does the solution provide mutual exclusion?

# Buggy #3

```
int turn = 0;
void lock (int id)  /*id = 0 or 1 */
{
    while(turn == id ^ 1);
}
void unlock (int id)
{
    turn = id ^ 1;
}
```

- Assuming  $T_0$  invokes lock( ) first, does the solution provide mutual exclusion?
- Yes it does, but there is another issue with this solution - two threads must request the lock in an alternate manner
- Progress requirement is not met
  - Argument: one of the threads stuck in an infinite loop



# Buggy #4

```
int flag[2] = {0,0}; int turn = 0;
void lock (int id)  /*id = 0 or 1 */
{
    turn = id ^ 1;
    flag[id] = 1;
    while(flag[id ^ 1] && turn == (id ^ 1));
}
void unlock (int id)
{
    flag[id] = 0;
}
```

- Why this solution does not work?

# Buggy #4

```
int flag[2] = {0,0}; int turn = 0;
void lock (int id) /*id = 0 or 1 */
{
    turn = id ^ 1;
    flag[id] = 1;
    while(flag[id ^ 1] && turn == (id ^ 1));
}
void unlock (int id)
{
    flag[id] = 0;
}
```

- Why this solution does not work?
- Mutual exclusion is not satisfied if  $T_0$  context switched after setting the turn = 1 and  $T_1$  acquires the lock (and sets turn = 0 in the process which allows  $T_0$  to acquire the lock)

# Attempt #5 (Peterson's solution)

```
int flag[2] = {0,0}; int turn = 0;
void lock (int id)  /*id = 0 or 1 */
{
    flag[id] = 1;
    turn = id ^ 1;
    while(flag[id ^ 1] && turn == (id ^ 1));
}
void unlock (int id)
{
    flag[id] = 0;
}
```

- Homework: Prove that mutual exclusion is guaranteed
- What about fairness?

# Attempt #5 (Peterson's solution)

```
int flag[2] = {0,0}; int turn = 0;
void lock (int id) /*id = 0 or 1 */
{
    flag[id] = 1;
    turn = id ^ 1;
    while(flag[id ^ 1] && turn == (id ^ 1));
}
void unlock (int id)
{
    flag[id] = 0;
}
```

- Homework: Prove that mutual exclusion is guaranteed
- What about fairness?
- The lock is fair because if two threads are contending, they acquire the lock in an alternate manner
- Extending the solution to N threads is possible

# Semaphores

- Mutual exclusion techniques allows exactly one thread to access the critical section which can be restrictive
- Consider a scenario when a finite array of size  $N$  is accessed from a set of producer and consumer threads. In this case,
  - At most  $N$  concurrent producers are allowed if array is empty
  - At most  $N$  concurrent consumers are allowed if array is full
  - If we use mutual exclusion techniques, only one producer or consumer is allowed at any point of time

# Operations on semaphore

```
struct semaphore{  
    int value;  
    spinlock_t *lock;  
    queue *waitQ;  
}sem_t;
```

// Operations

```
sem_init(sem_t *sem, int init_value);  
sem_wait(sem_t *sem);  
sem_post(sem_t *sem);
```

- Semaphores can be initialized by passing an initial value
- *sem\_wait* waits (if required) till the value becomes +ve and returns after decrementing the value
- *sem\_post* increments the value and wakes up a waiting context
- Other notations: P-V, down-up, wait-signal

# Unix semaphores

```
#include <semaphore.h>
```

```
main(){  
    sem_t s;  
    int K = 5;  
    sem_init(&s, 0, K);  
    sem_wait(&s);  
    sem_post(&s);  
}
```

- Can be used to in a multi-threaded process or across multiple processes
- If second argument is 0, the semaphore can be used from multiple threads
- Semaphores initialized with value = 1 (third argument) is called a binary semaphore and can be used to implement locks
- Initialize: `sem_init(s, 0, 1)`  
lock: `sem_wait(s)`, unlock: `sem_post(s)`

# Semaphore usage example: wait for child

```
child(){  
    ...  
    sem_post(s);  
    exit(0);  
}  
int main (void ){  
    sem_init(s, 0);  
    if(fork() == 0)  
        child();  
    sem_wait(s);  
}
```

- Assume that the semaphore is accessible from multiple processes, value initialized to zero
- If parent is scheduled after the child creation, it waits till child finishes
- If child is scheduled and exits before parent, parent does not wait for the semaphore



# Semaphore usage example: ordering

```
A=0; B=0;
```

```
Thread-0 {  
    A = 1;  
    printf("B = %d\n", B);  
}
```

```
Thread-1 {  
    B=1;  
    printf("A = %d\n", A);  
}
```

- What are the possible outputs?

# Semaphore usage example: ordering

A=0; B=0;

Thread-0 {

    A = 1;

    printf("B = %d\n", B);

}

- What are the possible outputs?
- (A = 1, B= 1), (A = 1, B = 0), (A = 0, B=1)
- How to guarantee A = 1, B= 1?

Thread-1 {

    B=1;

    printf("A = %d\n", A);

}

# Semaphore usage example: ordering

```
sem_init(&s1, 0);  
A=0; B=0;  
Thread - 0 {  
    A = 1;  
    sem_wait(&s1);  
    printf("B = %d\n", B);  
}  
Thread - 1 {  
    B=1;  
    sem_post(&s1);  
    printf("A = %d\n", A);  
}
```

- What are the possible outputs?

# Semaphore usage example: ordering

```
sem_init(&s1, 0);  
A=0; B=0;  
Thread - 0 {  
    A = 1;  
    sem_wait(&s1);  
    printf("B = %d\n", B);  
}  
Thread - 1 {  
    B=1;  
    sem_post(&s1);  
    printf("A = %d\n", A);  
}
```

- What are the possible outputs?
- (A = 1, B = 1), (A=0, B=1)
- How to guarantee A = 1, B= 1?

# Ordering with two semaphores

```
sem_init(s1, 0);  
sem_init(s2, 0);  
A=0; B=0;
```

- Waiting for each other guarantees desired output

Thread - 0

```
{  
    A = 1;  
    sem_post(s1);  
    sem_wait(s2);  
    printf("%d\n", B);  
}
```

Thread - 1

```
{  
    B=1;  
    sem_wait(s1);  
    sem_post(s2);  
    printf("%d\n", A);  
}
```